

CXR Dual Translate Mapping Curriculum
MAP-00...09 — frozen 2026-09-11 pass (mapping, not selection)

CXR evidence-grounding lab

Generated from markdown modules

Contents

How to use this PDF	3
MAP-00 — Orientation	4
Theory — say this out loud	4
GUI — two ports, do not mix them	4
Backend — trees	4
Week 1 — send-back	4
MAP-01 — What we set out to achieve	5
Theory — the original aim	5
Background — why Dual exists	5
What “several inputs” means here	5
Results you will need later	5
Backend — find the locked question	6
Week 1 — send-back	6
MAP-02 — Backend pipeline (extract, ground, rule)	7
Theory — three functions, three jobs	7
GUI	7
Backend — run this (CPU)	7
Files	8
Week 1 — send-back	8
MAP-03 — Three surfaces (why this pass is Translate)	9
Theory — Encode, Compute, Translate	9
Why “several approaches per note” still one class	9
GUI	9
Backend — locator is frozen	9
Week 1 — send-back	9
MAP-04 — Filling Dual extracts (Goal 2)	10
Theory — the empty-extract problem	10
GUI	10
Backend — inspect one frozen extract (no GPU)	10
Result (Goal 2)	10
Week 1 — send-back	11
MAP-05 — Nested Translate cells (G3-G9)	12
Theory — one class, six depths	12
Implementation (where to read)	12
Results after the stack (n=108 Dual-fill)	12
Week 1 — send-back	13
MAP-06 — Four notes walked (I2, C4, T1, T2)	14
Theory — how to read a ladder	14
Backend — one command	14
What you should see	14
GUI	14
Week 1 — send-back	14

MAP-07 — Leftovers and why we stopped	15
Theory — leftover is not a homework list	15
Results (Goal 10)	15
Stop rule (canon)	15
Week 1 — send-back	15
MAP-08 — Mapping is not selection (and the viewer)	16
Theory — two words people mix up	16
GUI — canonization workbench	16
Backend — meta API	16
Week 1 — send-back	16
MAP-09 — Replicate, and the claim ceiling	17
Replicate book (CPU, frozen artifacts)	17
Files that <i>are</i> the canon	17
Claim ceiling — say / never say	17
What would be a <i>new</i> investigation (not this PDF)	18
Week 1 — send-back	18

How to use this PDF

North-star: By MAP-09 you can explain Dual grounding, what this pass set out to do, how nested Translate patches work, how to replay them, and what the evidence does not support.

Not: a generic auto-corrector for every LLM failure. **Not** locator-chosen Encode vs Compute vs Translate. **Not** CXR Foundations or RepEng Ph9-14.

Role	Where
This PDF	Print / offline
Track UI	http://127.0.0.1:8264/
Map viewer	http://127.0.0.1:8263/
Science lab	cxr-evidence-grounding-lab/

Order: **MAP-00** → **MAP-09**.

MAP-00 — Orientation

Parent: CXR Dual Translate Mapping Curriculum

Next: [MAP-01](#)

Duration: 1 session

Theory — say this out loud

This track is about **one question on synthetic clinical notes**:

Did first-line therapy fail?

A language model reads the note and fills a JSON **extract**. A **fixed rule** (not the model) then says SATISFIED, NOT_SATISFIED, UNCERTAIN, or CONTRADICTION.

In 2026-09-11 we asked: when the Dual path is wrong, can we **group** those mistakes and **patch the extract** so the same rule becomes right — without flipping notes that were already right?

What we got: a nested map of Translate patches (G3 through G9), shown per note in a GUI.

What we did not get: a machine that automatically corrects every failure in a large language model.

Soft claim: mapping yes (partial). One-size-fits-all no. Selection among independent editors no. Auto-correct no.

GUI — two ports, do not mix them

Port	What it is
8264	This curriculum (reading)
8263	Map viewer (replay the frozen ladder on 108 notes)
8260	Frozen locator. Routing only. Do not expand it for this track

1. Open `http://127.0.0.1:8263/`
2. Read the yellow banner. It should say mapping, not selection.
3. Click **T1** (already correct) and **T2** (still miss). You will use those in MAP-06.

Start the viewer if it is down:

```
cd ~/staging/cxr-evidence-grounding-lab
./scripts/run_map_viewer.sh
```

Backend — trees

```
cd ~/staging/cxr-evidence-grounding-lab
ls n2s_lab/n2s_g{2,3,4,6,7,8,9,10}*.py n2s_lab/n2s_v1_selection_contrast.py n2s_lab/n2s_map_viewer.py
ls artifacts/n2s-g2-analog-fill-extracts.json artifacts/n2s-v1-selection-contrast.json
ls curriculum-mapping/modules
```

Call stack you will defend:

```
frozen Dual extract (Goal 2 JSON)
-> optional nested relabel G3..G9 (copy/wrap; not a new model)
-> ground() (unchanged)
-> evaluate_rule() (unchanged)
-> verdict vs gold
```

Week 1 — send-back

Three lines: (1) the Dual question, (2) mapping vs auto-correct, (3) which GUI is :8263.

MAP-01 — What we set out to achieve

Prerequisite: MAP-00 · **Next:** [MAP-02](#)

Theory — the original aim

The ambitious sentence was:

Can we automatically change and correct every failure mode inside a large language model?

That is a **later thesis**, not a closeable v1. It is open-ended: there is no finite list of “every failure.”

This pass asked a **smaller, empirical** question (design doc docs/TRACKB-FAILURE-INTERVENTION-MAP.md):

Can recurring Dual-shaped clinical-note failure signatures be mapped to correction methods that repair those failures while preserving unaffected notes — and could that map later support *choosing* an intervention?

Two halves:

1. **Mapping** — group similar Dual mistakes; attach a patch; check collateral (did we break notes that were already right?).
2. **Selection** — given a note, pick patch A vs B vs abstain, not “always apply the deepest stack.”

We completed a **partial mapping** on the Translate/extract side. We did **not** complete selection or auto-correct.

Background — why Dual exists

Doctors write messy notes. Rules need four crisp facts:

Atom	Plain meaning
A	A first-line (1L) therapy is identified
B	That therapy was actually given, not only planned
C	Something that counts as failure happened
D	That failure is <i>of</i> first-line, not some other line

Plus **X**: two claims in the note cannot both be true about those facts.

The model is allowed to fill the extract. The rule is not allowed to “use vibes.” If C is unknown, the rule says UNCERTAIN. If X is true, CONTRADICTION. If any of A–D is false, NOT_SATISFIED. Only A and B and C and D true, and not X, is SATISFIED.

That split is the **neural to symbolic boundary**. Fidelity means: did the meaning in the note survive into those atoms?

What “several inputs” means here

Each row is a **synthetic note** (not a real patient), with a gold label. Different notes do different clinical jobs:

- sequenced response then later progression (T1)
- failure implied by “platinum-refractory” (I2)
- two incompatible 1L stories on the same chart (C4)
- treatment just started, no failure yet (T2)

A one-size patch that helps I2 can hurt a P-neg note if you are careless. That is why we measured **P-neg**, **P-xspan**, and **P-bind** as arms, not only “miss count went down.”

Results you will need later

Original part	This-pass answer
Map Dual failures to corrections	Partial yes — six nested Translate cells, n=108

Original part	This-pass answer
Repair without breaking others	Yes on measured arms (P-neg leak 0 after G3; P-xspan 10/10 after G6)
Choose among independent editors	No — only one class (nested Translate)
Locator picks Encode vs Compute vs Translate	No
Auto-correct every LLM failure	No

Backend — find the locked question

```
cd ~/staging/cxr-evidence-grounding-lab
sed -n '1,12p' docs/TRACKB-FAILURE-INTERVENTION-MAP.md
```

Week 1 — send-back

Write the smaller Dual question in one sentence. Write the ambitious auto-correct sentence. Mark which one this curriculum canonizes.

MAP-02 — Backend pipeline (extract, ground, rule)

Prerequisite: MAP-01 · **Next:** [MAP-03](#)

This module is the **implementation** you must be able to replay without a GPU.

Theory — three functions, three jobs

```
note text
  | (language model, already run; we reuse frozen JSON)
  v
Extraction (quotes, admin status, outcome polarities, uncertainty, contradiction cues)
  | ground()
  v
Grounding (atoms A B C D, flag X)
  | evaluate_rule()
  v
Verdict    SATISFIED | NOT_SATISFIED | UNCERTAIN | CONTRADICTION
```

- `ground()` **does not call a model**. It is a deterministic map from extract fields to atoms. We **did not rewrite it** in this pass.
- `evaluate_rule()` is also frozen. Short version:

```
if X: CONTRADICTION
else if any of A,B,C,D is false: NOT_SATISFIED
else if any of A,B,C,D is unknown: UNCERTAIN
else: SATISFIED
```

Patches in this investigation edit the **extract** (Translate), then call the same `ground()` and the same rule.

GUI

On :8263, pick **T1**. The **Atoms** column is A B C D X after each cell. On T1 they never change: the raw extract already grounds to SATISFIED.

Backend — run this (CPU)

From the lab root. Uses the **frozen Goal 2 fill extracts**, not a live 7B call.

```
cd ~/staging/cxr-evidence-grounding-lab
../cxrlabs/faiss_gpu1/bin/python - <<'PY'
import json
from n2s_lab.n2s_g2_analog_fill import EXTRACT_PATH
from n2s_lab.ground import ground
from n2s_lab.predicate import evaluate_rule, PREDICATE_FORMULA
from n2s_lab.types import Extraction

print(PREDICATE_FORMULA)
raw = json.loads(EXTRACT_PATH.read_text())["extracts"]["T1"]["extract"]
ex = Extraction.from_dict(raw)
g = ground(ex)
r = evaluate_rule(g)
print("admin", ex.administration_status)
print("outcomes", [(o.text[:40], o.polarity) for o in ex.outcome_statements])
print("atoms", g.to_dict())
print("verdict", r.verdict.value, "rule", r.fired_rule)
PY
```

You should see T1 SATISFIED with A=B=C=D true and X false.

Read the rule text:

```
sed -n '14,45p' n2s_lab/predicate.py
```

Read that ground is the boundary and does not call a model:

```
sed -n '1,8p' n2s_lab/ground.py
```

Files

File	Role
n2s_lab/types.py	Extraction, Grounding, Verdict
n2s_lab/ground.py	extract → atoms
n2s_lab/predicate.py	atoms → verdict
n2s_lab/n2s_phase2_translate.py	atoms_payload() helper + later relabels
artifacts/n2s-g2-analog-fill-extracts.json	frozen Dual-fill extracts n=108

Week 1 — send-back

Paste your T1 verdict line. Name the three stages extract / ground / rule.

MAP-03 — Three surfaces (why this pass is Translate)

Prerequisite: MAP-02 · **Next:** [MAP-04](#)

Theory — Encode, Compute, Translate

A Dual miss can be *talked about* in three places. That is a **diagnostic partition**, not proof of where the error “first emerged.”

Surface	Rough question	What we had in this lab
Encode	Did the note land in the wrong representation?	L8 readout on two notes; no Encode editor
Compute	Did layer-to-layer computation lose the fact?	Lost-steps; frozen-d patch CLOSED ; G5 REVIEW = contain, not repair
Translate	Did the extract+ground drop meaning the model already wrote?	This is where the only demonstrated Dual repair lived

Phase 1-3 parked that story in docs/TRACKB-PHASE1-3-PORTFOLIO.md. Locator :8260 **routes** known notes to a surface label. It does **not** choose a repair class. `locator_chose_repair_class=false`.

This mapping pass therefore stayed on **Translate**: copy or wrap extract-relabel functions. We did not invent an Encode editor to fill a matrix. We did not reopen frozen-d.

Why “several approaches per note” still one class

On :8263 you see G3 then G4 then G6... That looks like several approaches. They are **nested**: G9 includes G8 includes ... includes G3. Choosing G9 is not choosing a different *kind* of intervention than G4. It is a **deeper stack of the same kind** (edit the extract, then `ground()`).

Independent classes would be things like: Translate patch vs hidden-state patch vs “leave it / REVIEW.” We only had the first.

GUI

Open :8260 only to remember it exists. Do not add cases. This curriculum’s live instrument is :8263.

Backend — locator is frozen

```
cd ~/staging/cxr-evidence-grounding-lab
sed -n '1,12p' docs/TRACKB-LOCATOR-V1.md
```

Do **not** run GPU frozen-d patches as part of this track.

Week 1 — send-back

One sentence: why mapping here is Translate-only. One sentence: why nested G3-G9 is not Encode vs Compute vs Translate selection.

MAP-04 — Filling Dual extracts (Goal 2)

Prerequisite: MAP-03 · **Next:** [MAP-05](#)

Theory — the empty-extract problem

Batch 0 Dual analog (user = note only) often produced **thin extracts**: almost no quotes, polarities unknown. Quote-promote then had nothing to copy. P-bind (failure sitting in quotes but not in outcome polarity) was **0/7**.

Goal 2 kept Dual analog (user=evidence) but put the extract **schema in the system prompt**. Same 7B, CPU/GPU as the original fill. Result: thin extracts **91 → 0** on n=108. That unblocked mapping. It is **not** Schema+note (a different prompt family we already knew changes extracts).

We **reuse** artifacts/n2s-g2-analog-fill-extracts.json. Do not overwrite it. Do not re-run the 7B fill unless you are deliberately reproducing Goal 2 (slow, GPU).

GUI

:8263 always starts from those Dual-fill extracts. Filter **Raw Dual-wrong**: 30 of 108 miss gold before any G3-G9 patch.

Backend — inspect one frozen extract (no GPU)

```
cd ~/staging/cxr-evidence-grounding-lab
../cxrlabs/faiss_gpu1/bin/python - <<'PY'
import json
from n2s_lab.n2s_g2_analog_fill import EXTRACT_PATH
from n2s_lab.types import Extraction

bundle = json.loads(EXTRACT_PATH.read_text())
print("n extracts", len(bundle["extracts"]))
rec = bundle["extracts"]["I2"]
print("parse_ok", rec["parse_ok"])
ex = Extraction.from_dict(rec["extract"])
print("admin", ex.administration_status)
print("quotes", ex.quotes)
print("notes", (ex.notes or "")[:200])
print("outcomes", [(o.polarity, o.text[:60]) for o in ex.outcome_statements])
PY
```

I2 often comes out administration_status=planned even though the note is platinum-refractory and considering a next line. That is the **shape** Goal 9 later patches. You are looking at the raw Dual fill, not G9 yet.

Optional live fill (only if you must reproduce Goal 2; not required for this curriculum):

```
# Slow. GPU. Do not overwrite if you only wanted to read.
# ./scripts/run_g2_analog_fill.sh
ls -l artifacts/n2s-g2-analog-fill-extracts.json
```

Result (Goal 2)

Metric	After Dual-fill
n	108
thin extracts	0 (was 91 on analog-empty)
Quote-promote still leaky	P-neg still fired until Goal 3 gated

Week 1 — send-back

Why analog-empty blocked mapping. What Goal 2 changed (schema in system, user still the note). I2 admin status you observed.

MAP-05 — Nested Translate cells (G3-G9)

Prerequisite: MAP-04 · **Next:** [MAP-06](#)

Theory — one class, six depths

Each “cell” is a **function** that takes an Extraction and returns a possibly edited Extraction. Then `ground()` and the rule run as usual.

Deeper functions **call the previous ones**. G9 = G8 + admin-blocked. G8 = G7 + same-time mix. And so on. There is no G5 repair cell (Goal 5 was a **census** only).

Cell	Plain job	What it typically edits
none	Raw Dual-fill extract	nothing
G3	Gated failure promote	asserted, non-negated failure words → failure polarity (not “no progression”)
G4	Mixed-span	simultaneous incompatible 1L claims → contradiction.present
G6	Never-quote	copy never / no-prior quotes into contradiction cues
G7	Restage	clear hedge once failure is confirmed; undo stable-then-progress X
G8	Same-time mix	undo notes-promote on uncertain; continue-now response → ongoing
G9	Admin-blocked	planned-after-refractory → given; fail vs never-dispensed cues

Do not treat these as six products you could ship independently. Nested means G6 still runs G3 and G4.

Implementation (where to read)

All live wrappers sit in `n2s_lab/n2s_phase2_translate.py` (copied/wrapped, not in-place mutation of G3 when G4 was added). The ladder order is declared in `n2s_lab/n2s_v1_selection_contrast.py` as `CELLS`.

```
cd ~/staging/cxr-evidence-grounding-lab
rg -n "^def relabel_failure_polarity" n2s_lab/n2s_phase2_translate.py
sed -n '50,58p' n2s_lab/n2s_v1_selection_contrast.py
```

Apply **one** depth the same way the GUI does:

```
../cxrlabs/faiss_gpu1/bin/python - <<'PY'
import json
from n2s_lab.n2s_g2_analog_fill import EXTRACT_PATH
from n2s_lab.n2s_phase2_translate import atoms_payload
from n2s_lab.n2s_v1_selection_contrast import _apply
from n2s_lab.types import Extraction

ex = Extraction.from_dict(
    json.loads(EXTRACT_PATH.read_text())["extracts"]["I2"]["extract"]
)
for name in ("none", "G3", "G9"):
    after, meta = _apply(ex, name)
    p = atoms_payload(after)
    print(name, p["verdict"], "admin", after.administration_status, "kind", meta.get("kind"))
PY
```

Expect I2: none/G3 still NOT_SATISFIED (planned); G9 SATISFIED (admin=given).

Results after the stack (n=108 Dual-fill)

After	Design P-bind match	n=108 Dual-wrong
raw (none)	—	30
G4	2/7 repaired path	27 miss
G6	—	19
G7	3/7	12
G8	3/7	8
G9	4/7 repaired; 7/7 match	6

P-neg leak **0** after G3. Design P-xspan **10/10** after G6. No new flips on G6-G9 vs the previous cell.

Panels (do not overwrite): artifacts/n2s-g3-gated-promote-panel.json ... n2s-g9-admin-blocked-panel.json.

Week 1 — send-back

Name the six cells. Say in one line why G9 is not a second editor class. Paste I2 none vs G9 verdicts from the snippet.

MAP-06 — Four notes walked (I2, C4, T1, T2)

Prerequisite: MAP-05 · **Next:** [MAP-07](#)

Theory without a note is fog. Run the same ladder the GUI uses.

Theory — how to read a ladder

For each depth: verdict, match vs gold, atoms, whether the extract actually changed.

- **Signature depth** (selector): phrase patterns in the extract that *would* pick that cell. Not a trained router.
- **Oracle**: shallowest depth that already matches gold. Upper bound, not a deployed selector.

If selector is none and oracle is none, the raw Dual extract was already right (T1).

If oracle is unrepaired, no cell in the stack matches gold (T2).

Backend — one command

```
cd ~/staging/cxr-evidence-grounding-lab
../cxrlabs/faiss_gpu1/bin/python curriculum-mapping/snippets/replay_ladder.py I2 C4 T1 T2
```

What you should see

I2 — implied failure, admin token wrong

Note: platinum-refractory, now considering next line. Gold **SATISFIED**.

- none: NOT_SATISFIED, admin=planned, B false
- G3: C and D become true (failure language), still planned so still miss
- G9: planned → given, B true, SATISFIED. Selector and oracle both G9

This is P-bind/P-impl: the clinic language was there; the extract parked “planned.”

C4 — two 1L stories at once

Gold **CONTRADICTION**.

- none/G3: SATISFIED (the rule believed a clean failure)
- G4: X true → CONTRADICTION. **Oracle = G4**
- G8: still CONTRADICTION, but may restage an outcome polarity. **Selector = G8** (deeper signature even though gold already matched)

This is why nested ≠ independent choice: G8 still includes G4.

T1 — already correct

Gold **SATISFIED**. Response then later progression. Raw extract already A=B=C=D true. Every cell no-op. Selector none, oracle none.

T2 — leftover (too early)

Gold **NOT_SATISFIED** (just started cycle 1; no failure event). Dual stays **UNCERTAIN**. No cell changes it. Oracle unrepaired. Tag S-too-early. Force-fitting a “failure” here would be the one-size-fits-all mistake.

GUI

On :8263, click the same four ids. Confirm the table matches the CLI (verdicts and “what this cell changed”).

Week 1 — send-back

Four lines: I2 / C4 / T1 / T2 — gold, first matching cell (or unrepaired), one-sentence why.

MAP-07 — Leftovers and why we stopped

Prerequisite: MAP-06 · **Next:** [MAP-08](#)

Theory — leftover is not a homework list

After G9, **6 / 108** Dual-fill notes still miss gold:

ID	Shape (Goal 10 tag)	Why not a new Translate cell
T2	too-early	one-off; failure has not happened
U2, U6	may-have	weak pair; BC_U2 already matches
TX_E03	held-for-progression	one-off
BC12_C1	missing pole / response-only mix	not a clean recurring phenotype
BC14_C1	naive $\neq$ never	one-off; do not stretch ground() to fit naive

Goal 10 tagged these shapes on **every** note (including already-correct ones). Recurrence was too weak to justify G11. `v1_freeze_recommended=true`.

Chasing miss→0 by writing an ID-fitted patch for T2 would **not** be mapping. It would be overfitting.

Results (Goal 10)

Artifact: `artifacts/n2s-g10-leftover-census.json`

```
cd ~/staging/cxr-evidence-grounding-lab
../cxrlabs/faiss_gpu1/bin/python - <<'PY'
import json
from pathlib import Path
p = json.loads(Path("artifacts/n2s-g10-leftover-census.json").read_text())
f = p["finding"]
print("n_miss", f["n_miss"])
print("ids_miss", f["ids_miss"])
print("new_cell justified?", f["shapes_that_justify_new_translate_cell"])
print("v1_freeze_recommended", f["v1_freeze_recommended"])
PY
```

If the JSON schema uses nested lists, open the file and find the six leftover ids. The map viewer filter **Still miss after G9** is the same six: T2, U2, TX_E03, U6, BC12_C1, BC14_C1.

Stop rule (canon)

- Do not add G11 leftover repairs.
- Do not rewrite ground() to absorb naive / too-early.
- Do not jump to CAA / frozen-d / Encode editor to “finish” Dual miss 6.
- The map is frozen: six nested cells + leftover census + viewer.

Week 1 — send-back

Why T2 should stay unrepaired. Why leftover 6 is a freeze, not a sprint backlog.

MAP-08 — Mapping is not selection (and the viewer)

Prerequisite: MAP-07 · **Next:** [MAP-09](#)

Theory — two words people mix up

Mapping: this extract *shape* is associated with this *kind of patch*, and we measured collateral.

Selection: at run time, choose among **independent** interventions (A vs B vs abstain), and show that the choice beats “always apply everything.”

We ran a contrast on the frozen stack (`n2s_lab/n2s_v1_selection_contrast.py`):

Arm	Miss / 108
none (raw Dual-fill)	30
extract-signature selector	6
always G9 (apply-all)	6
oracle (shallowest match)	upper bound; 79.6% agreement with selector

The selector chose G9 on **2 / 108** notes (not “always G9 in disguise”). Miss 6 = apply-all 6. That is **stack-depth routing from extract phrases**, not Encode vs Compute vs Translate, and not A vs B vs C vs abstain.

Walk-back (same day): nested wrappers are still **one class**. V1 demonstrated mapping, not class selection.

GUI — canonization workbench

`http://127.0.0.1:8263/`

```
cd ~/staging/cxr-evidence-grounding-lab
./scripts/run_map_viewer.sh
```

Use filters:

- **All** — 108 notes
- **Raw Dual-wrong** — 30
- **Still miss after G9** — 6 leftovers
- **Design P-bind** — the bind arm
- **Selector ≠ none** — notes whose extract signature is not none

Blue row = signature depth. Dashed outline = oracle. Blue bar = that cell changed the extract.

Frozen locator :8260 is a different product (routing). This viewer is not an expansion of it.

Backend — meta API

```
curl -sS http://127.0.0.1:8263/api/meta | python3 -c "import json,sys; m=json.load(sys.stdin); print(m['
```

Panel JSON: `artifacts/n2s-map-viewer-panel.json`

Contrast JSON: `artifacts/n2s-v1-selection-contrast.json`

Week 1 — send-back

In two sentences: what the selector measured, and why we still say “not selection.”

MAP-09 — Replicate, and the claim ceiling

Prerequisite: MAP-08

Unlock: you can teach MAP-00...08 without inflating auto-correct.

Replicate book (CPU, frozen artifacts)

From cxr-evidence-grounding-lab. Typical Python: `../cxrlabs/faiss_gpu1/bin/python`.

1. Self-tests (no GPU)

```
cd ~/staging/cxr-evidence-grounding-lab
../cxrlabs/faiss_gpu1/bin/python -m n2s_lab.n2s_map_viewer
../cxrlabs/faiss_gpu1/bin/python map_viewer_server.py --selftest
```

2. Ladder on four teaching notes

```
../cxrlabs/faiss_gpu1/bin/python curriculum-mapping/snippets/replay_ladder.py I2 C4 T1 T2
```

3. Viewer

```
./scripts/run_map_viewer.sh
# open http://127.0.0.1:8263/
```

4. Curriculum UI

```
./scripts/run_mapping_curriculum.sh
# open http://127.0.0.1:8264/
```

5. Optional: rebuild this PDF

```
./curriculum-mapping/modules/build-mapping-pdf.sh
```

Do not need, for this track: re-extracting 108 notes with 7B; Dual_full rescore; frozen-d; CAA; rewriting ground().

Files that are the canon

Path	Role
docs/TRACKB-FAILURE-INTERVENTION-MAP.md	Design + cell results
artifacts/n2s-g2-analog-fill-extracts.json	Frozen Dual-fill extracts
n2s_lab/n2s_phase2_translate.py	Nested relabels
n2s_lab/ground.py / predicate.py	Unchanged boundary
artifacts/n2s-g10-leftover-census.json	Why no G11
artifacts/n2s-v1-selection-contrast.json	Mapping vs apply-all
static/map-viewer.html	:8263
curriculum-mapping/	This track

Claim ceiling — say / never say

Say

- Dual-shaped extract failures can be grouped.
- Nested Translate relabels repaired a measured subset with controlled collateral on P-neg / P-xspan / P-bind.
- One-size-fits-all is the wrong default (T1 needs nothing; T2 should not be forced).
- This is a **partial positive** on mapping.

Never say

- We automatically correct every LLM failure mode.
- :8260 chose the right repair class.
- G3-G9 are six independent editors.
- Encode/Compute editors exist in this pass.
- Leftover 6 means the map is unfinished homework.

What would be a *new* investigation (not this PDF)

A second **independent** class (or principled abstain as a chosen non-repair), then A vs B vs abstain. Encode none today; Compute editor CLOSED. That is a new question. This curriculum stops.

Week 1 — send-back

List three commands you ran. Quote the honest claim in one sentence. Name one claim you will not make.